


Building a Quantum Computer with QLDPC Codes

Sunny Zhiyang He @ Simons Institute, July 23rd, 2026




Fault-Tolerant Quantum Computation in 2026

The shift from NISQ to FTQC

2018: “Because of the very hefty overhead cost of quantum error correction, **the era of fault-tolerant quantum computing (FQTC) may still be rather distant**. No one really knows how long it will take to get there.” — John Preskill, “Quantum Computing in the NISQ era and beyond”

2026: Industries converge to fault-tolerance, aim to deliver **hundreds to thousands of logical qubits, from megaquop to teraquop, before 2030**.



We went from ‘large-scale FTQC maybe in 20-30 years’ to ‘we might break crypto before 2030’.

What happened?

Convergence of experiments and QEC theory

QEC Theory

2019+: New theory for QLDPC codes, leading to asymptotically good codes in 2021

2019+: New decoding algorithms, focus on BP and ML

2020+: Development of simulation software, such as STIM and better decoders

2023+: Focus on finite-size QLDPC codes and their implementation on hardware

2023+: Computational methods for LDPC codes, including QLDPC surgery and transversal logic

Experiments

2022+: Reconfigurable qubits enable non-local connectivity

2023+: Sub-threshold surface code memory, larger code gives more error suppression

2023+: Long-range connection that supports modular architectures

2024+: Magic state distillation and cultivation, gate teleportation

Consistently: Lower gate error rate, larger systems, better physical control

Finite-scale proposals to reduce space overhead, with focus on LDPC codes

New hardware capabilities, lower physical and logical error rates, larger scale systems

Result of fusion: new proposals of finite-size FTQC

Gidney 2025 estimate [2505.15917]:

- Surface code architecture with lattice surgery, applied magic state cultivation and code concatenation.
- RSA-2048 in less than 1 million qubits, less than 1 week, on superconducting qubits.

IBM Bicycle Architecture [2506.03094]:

- Modular LDPC architecture, computation by LDPC surgery.
- Roughly 10 times less qubits and 10 times slower than surface codes.

Iceberg [2602.11457]:

- Modular LDPC architecture, computation by LDPC surgery.
- RSA-2048 in 100000 qubits, 30000 days, on neutral atoms.

Oratomic [2603.28627]:

- Modular LDPC architecture with high-rate codes, computation by LDPC surgery.
- ECC-256 in 10000 qubits, 1000 days, on neutral atoms. RSA-2048 in 11000 qubits, 43000 days.

More recent proposals: IonQ Walking Cat [2604.19481], Duke + UT Austin + Yale [2604.19735], QuEra's LDPC STAR [2606.25011], Google's modular hyperbolic architecture [presented at QEC26], and more.

Caution

These are proposals and estimates, not finalized reports. Some assumptions and methods are under discussion in the community.

Result of fusion: new proposals of finite-size FTQC

Gidney 2025 estimate [2505.15917]:

- Surface code architecture with lattice surgery, applied magic state cultivation and code concatenation.
- RSA-2048 in less than 1 million qubits, less than 1 week, on superconducting qubits.

IBM Bicycle Architecture [2506.03094]:

- **Modular LDPC architecture**, computation by LDPC surgery.
- Roughly 10 times less qubits and 10 times slower than surface codes.

Iceberg [2602.11457]:

- **Modular LDPC architecture**, computation by LDPC surgery.
- RSA-2048 in 100000 qubits, 30000 days, on neutral atoms.

Oratomic [2603.28627]:

- **Modular LDPC architecture** with high-rate codes, computation by LDPC surgery.
- ECC-256 in 10000 qubits, 1000 days, on neutral atoms. RSA-2048 in 11000 qubits, 43000 days.

More recent proposals: IonQ Walking Cat [2604.19481], Duke + UT Austin + Yale [2604.19735], QuEra's LDPC STAR [2606.25011], Google's modular hyperbolic architecture [presented at QEC], and more.

Observation 1

All (except Gidney's) architectures are done with LDPC codes. We see **dramatic reduction in space overhead**, which has been the key bottleneck for FTQC.

Result of fusion: new proposals of finite-size FTQC

Gidney 2025 estimate [2505.15917]:

- Surface code architecture with lattice surgery, applied magic state cultivation and code concatenation.
- RSA-2048 in less than 1 million qubits, less than 1 week, on superconducting qubits.

IBM Bicycle Architecture [2506.03094]:

- Modular LDPC architecture, computation by LDPC surgery.
- Roughly 10 times less qubits and 10 times slower than surface codes.

Iceberg [2602.11457]:

- Modular LDPC architecture, computation by LDPC surgery.
- RSA-2048 in 100000 qubits, 30000 days, on neutral atoms.

Oratomic [2603.28627]:

- Modular LDPC architecture with high-rate codes, computation by LDPC surgery.
- ECC-256 in 10000 qubits, 1000 days, on neutral atoms. RSA-2048 in 11000 qubits, 43000 days.

More recent proposals: IonQ Walking Cat [2604.19481], Duke + UT Austin + Yale [2604.19735],

QuEra's LDPC STAR [2606.25011], Google's modular hyperbolic architecture [presented at QEC], and more.

Observation 2

LDPC surgery serves as the computational backbone for half of these proposals.

Result of fusion: new proposals of finite-size FTQC

Gidney 2025 estimate [2505.15917]:

- Surface code architecture with lattice surgery, applied magic state cultivation and code concatenation.
- RSA-2048 in less than 1 million qubits, **less than 1 week**, on superconducting qubits.

IBM Bicycle Architecture [2506.03094]:

- Modular LDPC architecture, computation by LDPC surgery.
- Roughly 10 times less qubits and **10 times slower than surface codes**.

Iceberg [2602.11457]:

- Modular LDPC architecture, computation by LDPC surgery.
- RSA-2048 in 100000 qubits, **30000 days**, on neutral atoms.

Oratomic [2603.28627]:

- Modular LDPC architecture with high-rate codes, computation by LDPC surgery.
- ECC-256 in 10000 qubits, 1000 days, on neutral atoms. RSA-2048 in 11000 qubits, **43000 days**.

More recent proposals: IonQ Walking Cat [2604.19481], Duke + UT Austin + Yale [2604.19735], QuEra's LDPC STAR [2606.25011], Google's modular hyperbolic architecture [presented at QEC], and more.

Observation 3

Significant increase in time cost,
to the point of infeasible.

Is this fundamental?


❖ No, not yet.

Outline

- I. Fault-Tolerant Quantum Computation in 2026
- II. Modern QLDPC Architectures
- III. Fault-Tolerant Computation by Surgery
- IV. Outlook

Modern QLDPC Architectures

Sketch of fault-tolerant architectures




Memory with QLDPC codes

Quantum low-density parity-check (LDPC) codes have **high encoding rate, at the cost of higher connectivity** requirement. Examples:

- Bivariate Bicycle code [144, 12, 12]^[1]
- Lifted/balanced product code [544, 80, ≤ 12]^[2]
- APM Codes [1152, 580, ≤ 12]^[3]

Surface code: [144, 1, 12]. **Space overhead of encoding can be reduced by more than 100x!**



IBM's bivariate bicycle code

Common memory designs:

- Few big blocks
- Many small blocks
- Concatenated codes^[4,5]

Tradeoff among modularity, space-efficiency, and access speed.

$$H_X = \begin{pmatrix} F_0 & F_1 & F_2 & F_3 & F_4 & F_5 & G_0 & G_1 & G_2 & G_3 & G_4 & G_5 \\ F_5 & F_0 & F_1 & F_2 & F_3 & F_4 & G_5 & G_0 & G_1 & G_2 & G_3 & G_4 \\ F_4 & F_5 & F_0 & F_1 & F_2 & F_3 & G_4 & G_5 & G_0 & G_1 & G_2 & G_3 \end{pmatrix}$$

$$H_Z = \begin{pmatrix} G_0^T & G_5^T & G_4^T & G_3^T & G_2^T & G_1^T & F_0^T & F_5^T & F_4^T & F_3^T & F_2^T & F_1^T \\ G_1^T & G_0^T & G_5^T & G_4^T & G_3^T & G_2^T & F_1^T & F_0^T & F_5^T & F_4^T & F_3^T & F_2^T \\ G_2^T & G_1^T & G_0^T & G_5^T & G_4^T & G_3^T & F_2^T & F_1^T & F_0^T & F_5^T & F_4^T & F_3^T \end{pmatrix}$$

APM codes of rate 1/2

¹ [IBM 2308.07915], ² [Xu et al. 2308.08648], ³ [Zhao et al. 2604.16209], [Kasai 2601.08824] ⁴ [Gidney et al. 2312.04522], ⁵ [Wills et al. 2605.21898].



Resource state factories

Almost all FTQC proposals use resource states, notably magic states for T and CCZ gates.

Recent architectures combine two primitives:

- **Magic state cultivation:** prepare a good ($10^{-5\sim 9}$) magic state at low distance, rapidly escape into large distance surface code.
- **Magic state distillation:** take many magic states at large distance, distill them into fewer states of better quality (10^{-12} and less).
- Typical factory: **cultivate first in surface code, load into LDPC codes and distill.**

Magic state production is one of the key bottlenecks in current QLDPC architectures! More on that later.



Computation machinery

For LDPC codes, we have two main methods of computation:


low-depth unitaries and **Pauli measurements**.

Low-depth unitaries:


- Typical example: transversal CNOT
- Constant-depth physical circuits that implement a logical unitary
 - + **Parallel, fast, no space overhead**
 - **Logical control is constrained and often limited**
 - **Usually needs qubit reconfiguration**

Pauli measurements:

- Usually implemented by code surgery, sometimes by CAT states
- Compute by extracting Pauli observables from logical qubits.
 - + **Complete logical control, can work with fixed connectivity**
 - **Usually slower, logical control is sequential in many cases**



Transversal CNOT on surface code^[1]



Lattice surgery^[2]

Why surgery now?

LDPC surgery^[1-6] is a new technique that has advanced rapidly.


Modern surgery techniques has many desirable properties:

- Provably fault-tolerant and low-overhead^[3,4]
- Complete logical control (often sequential)^[6]
- Effective starting at small scale: <400 physical qubits enable control of 11 logical qubits at distance 12^[2,7]
- Implementable on different qubit platforms, including fixed-connectivity hardware^[4,6]

Surgery and low-depth unitaries are similar to CPU and GPU:

sequential and complex control, vs. parallel and structured control.

- In early FTQC, as we want to implement specific tasks ASAP, surgery is a very natural choice.



Extractor architecture for
LDPC surgery-based FTQC^[6]

¹[Cohen et al. 2110.10794], ²[Cross et al. 2407.18393], ³[Williamson & Yoder 2410.02213], ⁴[Ide et al. 2410.02753], ⁵[Swaroop et al. 2410.03628],

⁶[He et al. 2503.10390]. ⁷[IBM 2506.03094]

If FTQC is like printing...

Suppose you have two different printing presses:

- One that prints 5 words at a time, any words of your choice;
- Another that prints 100 copies of one chosen word at a time;


When resource is limited and you really want to print one specific essay, use the first press

- This is **sequential surgery**, which was studied in [1, 2] and adopted in recent QLDPC architectures;

When you want to print many copies of the same essay, use the second press

- This is **batched, parallel computation via transversal gates**^[3]

* Surgery can give parallel, structured logical control^[4] as well, more on that later.



Fault-Tolerant Computation by Surgery

Universal computation via logical measurements


A Clifford + T circuit can be compiled into two primitives:

1. measurement of Pauli observables
(Pauli-product measurements, PPMs)
2. magic states production

This is called **Pauli-based computation**.


Code surgery is a way to measure logical Pauli observables fault-tolerantly.

- Starting at a quantum code (a stabilizer group), we measure a new stabilizer group, where the (large) logical Pauli become a product of (small) new stabilizers.
- Weight reduction → fault-tolerance



Surface code lattice surgery

Logical measurements on surface codes: **lattice surgery**, [Horsman et al, 1111.4022].



Product of **red** X-checks = $X_L \otimes X_L$ – **obtain logical measurement result by measuring new stabilizers.**

LDPC surgery via Extractors


Logical control on high-rate ($k = O(n)$) codes is fundamentally different from that on constant-dimension ($k = O(1)$) codes. For PPMs, there are 4^k many logical Pauli operators.

Current model of sequential surgery:


For *any* code, given *any* set of logical Paulis, can build an **extractor system** that enable measurement of *any* operator from the set, *one at a time*.

➤ Extractor size ranges between $\tilde{O}(d) \sim \tilde{O}(n)$.

Key idea: extractors augment QLDPC memories into processors.




An extractor-augmented computational (EAC) block.



Pauli measurement with extractors


On a high level: for a logical operator $\mathcal{L}$ on a code Q , define a new stabilizer group supported on Q and $\mathcal{X}$. Measure the new stabilizer group (for d rounds) to measure $\mathcal{L}$.



Extractor architectures for Pauli-based FTQC

Building an architecture with extractors is like putting together Lego pieces:


1. Choose your QLDPC codes/memory
2. Choose your logical control (PPMs) and extractor systems
 - This choice deeply impacts/depends on compilation method and resource budget!
3. Choose your magic state factories
4. Connect them up with adapters^[1,2]



Design choices of recent architectures


Similarity: **sequential logic**

- Every CCZ (or T) gate in the circuit is compiled into one (large) PPM. Many (or all) Clifford gates are compiled away.
- Magic state throughput is **one/logical cycle**.



Differences:

- Choice of codes – heavily impact space overhead
- Choice of extractors, including loading – heavily impacts time overhead
- Choice of magic state factories – cost of factory limits magic throughput
- Physical assumptions on qubits
- Compilation and resource estimation target
- Many more details



Oratomic's architecture^[3]



- MEMORY (75 kq)
- PROCESSING UNIT (15 kq)
- MAGIC ENGINE (5 kq)

Iceberg's architecture^[2]

Demystifying the numbers

Back-of-envelop space estimates for leading factors

Memory: Use high-rate codes^[1-3], ~1500 logical qubits in 5000~9000 physical qubits.


Compute: Can use smaller processor codes, or compute directly on memory.

- For simplicity, we can use single-operator extractors to execute large PPMs directly on memory.
- Rule of thumb: roughly the same size as memory.

Magic states: cultivate to d~15 surface codes, load into n~500 QLDPC codes using extractors, distill.

- Many designs ranging from 5000~35000 qubits.
- Need detailed simulation to pin down exact costs.

Total: 15k ~ 55k physical qubits.



Baseline: [Gidney 2025 estimate \[2505.15917\]](#)

- Surface code architecture with lattice surgery
- RSA-2048 in less than 1 million qubits, less than 1 week, on superconducting qubits
- One CCZ state per logical cycle, some parallelism in Clifford operations.

¹[Kasai, 2601.08824], ²[Chen et al. 2604.16209], ³[Oratomic 2603.28627].

Demystifying the numbers

Back-of-envelop time estimates for leading factors

Hardware: neutral atom code cycle 1 *ms* vs.
superconducting code cycle 1 μ s. Factor of 1000!

- Takes Gidney's 5 days to 5000 days.


Compilation:

- Load in/out between memory and compute;
- Compilation of PPMs;
- Choice of magic states (T vs. CCZ);
- Many more factors which depends on architectural and algorithmic choices.

Logical clock speed:

- In current estimates, 1 logical cycle = $O(d) \sim 20$ physical cycles. Varies with assumptions.

Total: on the order of 10k days.



Baseline: [Gidney 2025 estimate \[2505.15917\]](#)

- Surface code architecture with lattice surgery
- RSA-2048 in less than 1 million qubits, less than 1 week, on superconducting qubits
- One CCZ state per logical cycle, some parallelism in Clifford operations.

Time overhead is not fundamental (yet)

Magic state throughput:

- Can we do more than 1 CCZ per logical cycle?
- Can we consume more than 1 CCZ per logical cycle?



Magic state production is the key bottleneck right now!

- Can we design more efficient factories?
- Parallel extractors can be built for consumption.^[1-3]

Hardware: can reconfigurable qubits move faster?



Need to ask our experimental overlords. However, we can implement surgery on superconducting qubits, which are a lot faster!

Logical cycle time:

- In current estimates, 1 logical cycle = $O(d) \sim 20$ physical cycles
- Can this be reduced to $O(1)$?



Yes! There are ways to **implement surgery with constant time overhead**.^[4]

Compilation:

- If the extractor is limited, a large PPM would need to be compiled into supported 'native instructions'
- **Leads to the 10x slow down in bicycle architecture**



Can be avoided by building a bigger extractor.

Challenges

Building a LDPC quantum computer will not be easy! **Many challenges remain, and we need lots of solid work to resolve them.**

Decoding:

- LDPC decoding is *much harder* than surface code decoding.
- Amazing progress in recent years, but we still need lots of work to improve speed *and* accuracy.
- **Good vs. bad decoding can lead to order-of-magnitudes difference in logical performance.**

Hardware complexity:

- Implementing QLDPC codes & surgery won't be easy, even for reconfigurable qubits.
- **For superconducting qubits, extractors can enable universal FTQC.** However, the connectivity needed for ultra-high-rate codes (rate $1/2$ or $1/3$) is more than current assumptions.

More QEC theory:

- To enable large-scale FTQC, we need further reduction in space-time overheads.
- Many concrete problems here.

Future directions for QLDPC architectures

Design more efficient magic state factories

- Can we find practical codes with low-depth CCZ (or other non-Clifford) gates?^[1,2]

Structured logical control co-designed with algorithm


- Algorithms and their subroutines often structured. Can we design the logical control (surgery or otherwise) with respect to their structure?^[3]

Practical design of (parallel/fast/full) extractors

- Surgery is very versatile, lots of opportunity for practical/finite-scale studies.

Is time-overhead fundamental?

- Does more efficient encoding of logical information necessarily lead to slowdown in logical computation (due to addressability or routing costs)?^[4]



Outlook

Long story short...

Experimental advance and QLDPC theory led to drastic reduction in space overhead for FTQC.

- The timeline advanced substantially. We will get early FTQC soon.
- The crucial bottleneck right now is time overhead.
- QEC will continue to improve, we need to switch to post-quantum cryptography.

Extractors and low-depth unitaries are analogous to classical CPU and GPU

- For the next phase of FTQC, a key aspect is joint-optimization of algorithms and QEC.
- FTQC designs will continue to evolve as we develop new methods, reach new hardware capabilities, and reach for new algorithmic goals.



Craig Gidney

@craiggidney.bsky.social

I would bet against Q day by 2030, but I wouldn't bet against it at 10:1 odds. ~10% risk is unacceptably high here, so I'm very in favor of transitioning to quantum-safe cryptography by 2029: [blog.google/innovation-a...](#)

Yes this means I 90% expect to be made fun of in 2030. Oh well.



Quantum frontiers may be closer than they appear

An overview of how Google is accelerating its timeline for post-quantum cryptography migration.

© blog.google

Quantum applications beyond factoring

Once the first generation of FTQC are built, we will be looking for experiments to run on them.


This would be the perfect time to test out new candidates for quantum applications!

We'd love to hear your ideas and (crazy) proposals!

- Worst-case proof of exponential time advantage is always nice but not required.
- Large polynomial speedup sounds awesome. 
- The goal is to discover practically/scientifically useful applications. It's ok if their performance at the near-term FTQC scale can be classically approximated.
- # of logical qubits is not as constrained now (hundreds of qubits is reasonable, thousands may need to wait a bit longer). 
- # of non-Clifford gates (T, CCZ, or other primitive) is (still) a good benchmark.
- From a FTQC perspective, we prefer circuits with more structure, such as parallel subroutines.

Building a Quantum Computer with QLDPC Codes

Sunny Zhiyang He @ Simons Institute, July 23rd, 2026



Slides

